

Other Retrieval Methods

Simon Pfreundschuh

December 19, 2017

Overview

1. Background
2. Markov Chain Monte Carlo
3. Bayesian Monte Carlo Integration
4. Machine Learning

Background

- ▶ We still want to solve the retrieval problem:

$$p(\mathbf{x}|\mathbf{y}) \propto p(\mathbf{y}|\mathbf{x})p(\mathbf{x}) \quad (1)$$

- ▶ Rodgers tells use how to do that if
 - ▶ $p(\mathbf{x}), p(\mathbf{y}|\mathbf{x})$ is Gaussian.
 - ▶ We have a(n at most *moderately non-linear*) forward model F
- ▶ But practically F should also
 - ▶ provide Jacobians,
 - ▶ not be too computationally complex.

Background

The Solutions

- ▶ Markov Chain Monte Carlo:
 - ▶ If you have a forward model and need the *full posterior distribution*.
- ▶ Database Retrievals:
 - ▶ (Bayesian) Monte Carlo Integration (BMCI)
 - ▶ Machine Learning
 - ▶ If performance is critical or if you don't have a forward model.

Markov Chain Monte Carlo

Advantages

- ▶ Direct sampling of the a posteriori distribution $p(\mathbf{x}|\mathbf{y})$
- ▶ No Jacobians required
- ▶ Not based on any assumptions (except a priori, of course)

Disadvantages

- ▶ Very Slow

Markov Chain Monte Carlo

- ▶ General method to sample from arbitrary distributions
- ▶ Sequential sampling: Next sample depends only on current state
 - ▶ *Markov chain property*
- ▶ Samples converge to target distribution


Basic Idea

- ▶ Sample *proposal state* from a proposal distribution (often a random walk).
- ▶ Accept or discard proposal depending on change in likelihood
- ▶ Repeat until convergence

Markov Chain Monte Carlo

A Retrieval Example

- ▶ Retrieval of integrated column water vapor from passive microwave observations
- ▶ 5 Channels: 23.8 GHz, 88.2 GHz, 165.5 GHz, 2×183.3 GHz
- ▶ Retrieve temperature and water vapor profiles



Markov Chain Monte Carlo

Let $p(\mathbf{x}|\mathbf{y})$ be the posterior we want to sample from and $J(\mathbf{x}_j|\mathbf{x}_i)$ a *symmetric proposal distribution*, i.e. satisfying $J_t(\mathbf{x}_j|\mathbf{x}_i) = J(\mathbf{x}_i|\mathbf{x}_j)$.

Metropolis Algorithm

1. Draw a starting point \mathbf{x}_0 with $p(\mathbf{x}_0|\mathbf{y}) > 0$
2. Iterate for $t = 1, \dots, n$:
 - 2.1 Sample a proposal \mathbf{x}^* from the proposal distribution $J_t(\mathbf{x}|\mathbf{x}_{t-1})$.
 - 2.2 Calculate

$$r = \frac{p(\mathbf{x}^*|\mathbf{y})}{p(\mathbf{x}_{t-1}|\mathbf{y})} \quad (2)$$

2.3 Set

$$\mathbf{x}_t = \begin{cases} \mathbf{x}^* & \text{with probability } \min(r, 1) \\ \mathbf{x}_{t-1} & \text{otherwise.} \end{cases} \quad (3)$$

Markov Chain Monte Carlo

Why does it work?

A Markov chain is guaranteed to have a *unique stationary distribution* if

- ▶ it is *aperiodic*,
- ▶ not *transient*,
 - ▶ i.e. there is no state that is not recurrent,
- ▶ *irreducible*,
 - ▶ i.e. there is no state for which there is a non-reachable state.

Thus only need to show that the stationary distribution is the posterior $p(\mathbf{x}|\mathbf{y})$.

Markov Chain Monte Carlo

Why does it work?

- ▶ Assume $p(\mathbf{x}_t|\mathbf{y}) = p(\mathbf{x}|\mathbf{y})$, i.e. the true posterior
- ▶ Consider the probability for a transition from \mathbf{x}_{t+1} to \mathbf{x}_t :

$$p(\mathbf{x}_{t+1}, \mathbf{x}_t | \mathbf{y}) = p(\mathbf{x}_t | \mathbf{y}) J(\mathbf{x}_{t+1} | \mathbf{x}_t) \min\left(\frac{p(\mathbf{x}_{t+1} | \mathbf{y})}{p(\mathbf{x}_t | \mathbf{y})}, 1\right) \quad (4)$$

$$= \operatorname{argmax}_{\mathbf{x}_t, \mathbf{x}_{t+1}} \{p(\mathbf{x} | \mathbf{y})\} J(\mathbf{x}_{t+1} | \mathbf{x}_t) \quad (5)$$

- ▶ This is symmetric as well: $p(\mathbf{a}, \mathbf{b}) = p(\mathbf{b}, \mathbf{a})$
- ▶ Symmetry of the joint distribution implies equality of the marginal distributions:

$$p(\mathbf{x}_{t+1} | \mathbf{y}) = p(\mathbf{x}_t | \mathbf{y}) = p(\mathbf{x} | \mathbf{y}) \quad (6)$$

Markov Chain Monte Carlo

Things to Consider

- ▶ If the posterior probability of a proposed state is higher than that of the current state, the proposal is always accepted.
 - ▶ The state will move towards high posterior densities.
- ▶ Algorithm needs time to reach stationary distribution.
 - ▶ Samples from *warm up* phase must be discarded.
- ▶ Consecutive samples are not independent.
 - ▶ Keep only every n th sample

Bayesian Monte Carlo Integration (BMCI)

- ▶ MCMC is (conceptually) nice, but also inherently slow.
- ▶ BMCI uses a database of *precomputed simulations* or *observations*.

General Idea

- ▶ Use a database of pairs (\mathbf{y}, \mathbf{x}) of observations and known \mathbf{x}
- ▶ Use importance sampling to transform samples in database to samples of the posterior

Bayesian Monte Carlo Integration (BMCI)

Consider the expected value $\mathcal{E}_{\mathbf{x}|\mathbf{y}}\{f(\mathbf{x})\}$ of a function f computed with respect to the a posteriori distribution $p(\mathbf{x}|\mathbf{y})$:

$$\int f(\mathbf{x}')p(\mathbf{x}'|\mathbf{y}) d\mathbf{x}' \quad (7)$$

Using Bayes theorem, the integral can be computed as

$$\int f(\mathbf{x}')p(\mathbf{x}'|\mathbf{y}) d\mathbf{x}' = \int f(\mathbf{x}') \frac{p(\mathbf{y}|\mathbf{x}')p(\mathbf{x}')}{\int p(\mathbf{y}|\mathbf{x}'') d\mathbf{x}''} d\mathbf{x}' \quad (8)$$

$$= \int f(\mathbf{x}')w(\mathbf{y}, \mathbf{x}')p(\mathbf{x}') d\mathbf{x}' \quad (9)$$

$$= \mathcal{E}_{\mathbf{x}}\{f(\mathbf{x})w(\mathbf{y}, \mathbf{x})\} \quad (10)$$

Bayesian Monte Carlo Integration (BMCI)

If the database is distributed according to our a priori assumptions, we can thus approximate any integral over the posterior distribution by:

$$\int f(\mathbf{x}') p(\mathbf{x}'|\mathbf{y}) d\mathbf{x}' \approx \sum_{i=1}^n f(\mathbf{x}_i) w(\mathbf{y}, \mathbf{x}_i) \quad (11)$$

Bayesian Monte Carlo Integration (BMCI)

The Weighting Function

- Assuming the database is exact up to a zero-mean, Gaussian error with covariance matrix \mathbf{S}_e the weighting function $w(\mathbf{y}, \mathbf{x})$ is given by:

$$w(\mathbf{y}, \mathbf{x}_i) = \frac{1}{C} \cdot \exp \left\{ -\frac{(\mathbf{y} - \mathbf{y}_i)^T \mathbf{S}_e^{-1} (\mathbf{y} - \mathbf{y}_i)}{2} \right\} \quad (12)$$

with normalization factor C

$$C = \int w(\mathbf{y}, \mathbf{x}) d\mathbf{x} \approx \sum_{i=1}^n w(\mathbf{y}, \mathbf{x}_i) \quad (13)$$

Bayesian Monte Carlo Integration (BMCI)

The Retrieval

- This can be used to retrieve the mean and variance of the posterior distribution:

$$\bar{x} = \mathcal{E}_{x|y}\{x\} \approx \sum_{i=1}^n w(\mathbf{y}, x_i) x_i \quad (14)$$

$$\text{var}(x) = \mathcal{E}_{x|y}\{(x - \bar{x})^2\} \approx \sum_{i=1}^n w(\mathbf{y}, x_i) (x_i - \mathcal{E}_{x|y}\{x\})^2 \quad (15)$$

- Or even the CDF of the posterior:

$$F_{x|y}(\mathbf{x}) = \int_{-\infty}^{\mathbf{x}} p(\mathbf{x}') d\mathbf{x}' \quad (16)$$

$$\approx \sum_{\mathbf{x}_i < \mathbf{x}} w(\mathbf{y}, \mathbf{x}_i) \quad (17)$$

Bayesian Monte Carlo Integration (BMCI)

Things to Consider

- ▶ A very large database may be required to truthfully represent the a priori and provide sufficient a posteriori statistics.
 - ▶ Solution: Weighting/clustering of database samples
- ▶ Traversing the database can take quite some time.
 - ▶ Solution: Sorting the database in a smart way

Bayesian Monte Carlo Integration (BMCI)

Example: Global Precipitation Measurement (GPM) Retrieval

International satellite mission to provide next-generation observations of rain and snow worldwide every three hours.

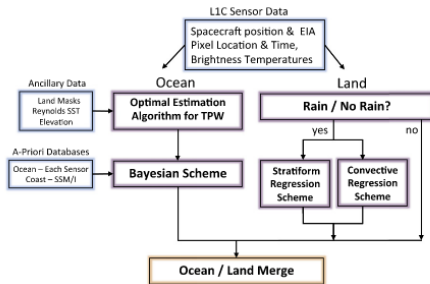


Figure: GPROF 2010 Retrieval Algorithm Flow (Kummerow et al. 2015).

Bayesian Monte Carlo Integration (BMCI)

Example: Global Precipitation Measurement (GMP) Retrieval

- ▶ Over Ocean:
 - ▶ Uses simulated database generated from profiles observed by the TRMM precipitation radar
 - ▶ Input: TBs, total precipitable water (TPR) from OEM, sea surface temperature (SST) from NWP
 - ▶ database with 65×10^6 entries stratified into SST/TPR bins of width 1 K / 1 mm.
 - ▶ Clustering algorithm used on bins to improve retrieval speed

Bayesian Monte Carlo Integration (BMCI)

Example: Global Precipitation Measurement (GMP) Retrieval




Figure: Trends in oceanic precipitation (Kummerow et al. 2015).

Bayesian Monte Carlo Integration (BMCI)

Example: Global Precipitation Measurement (GMP) Retrieval

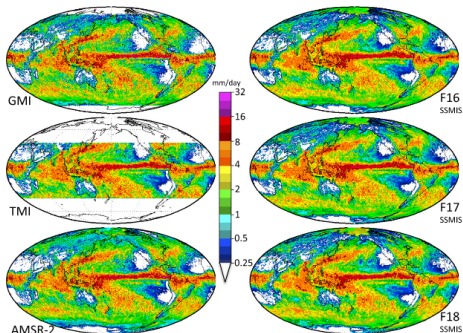


FIG. 10. April–June 2014 average mean surface precipitation of the GPM constellation conical sensors.

Figure: Precip from GPROF (Kummerow et al. 2015).

Machine Learning

Idea

- ▶ Try to learn the inverse method $\mathbf{x} = R(\mathbf{y})$ directly from the data.
- ▶ Regression is an old problem: Plenty of methods to choose from
 - ▶ Traditional regression analysis
 - ▶ Machine Learning

Machine Learning

Neural Networks

- Universal estimators that compute a vector of output activations $\mathbf{y} = F_{NN}(\mathbf{x}_i, \mathbf{W}_i, \boldsymbol{\theta}_i)$ from a vector of input activations \mathbf{x} :

$$\mathbf{x}_0 = \mathbf{x} \tag{18}$$


$$\mathbf{x}_i = f_i(\mathbf{W}_i \mathbf{x}_{i-1} + \boldsymbol{\theta}_i) \tag{19}$$

$$\mathbf{y} = \mathbf{x}_n \tag{20}$$

- Weight matrices \mathbf{W}_i and bias vectors $\boldsymbol{\theta}_i$ are *learnable parameters* of the network.


Machine Learning

Neural Networks




Machine Learning

Neural Networks




Machine Learning

Neural Networks




Machine Learning

Neural Networks




Machine Learning

Neural Networks



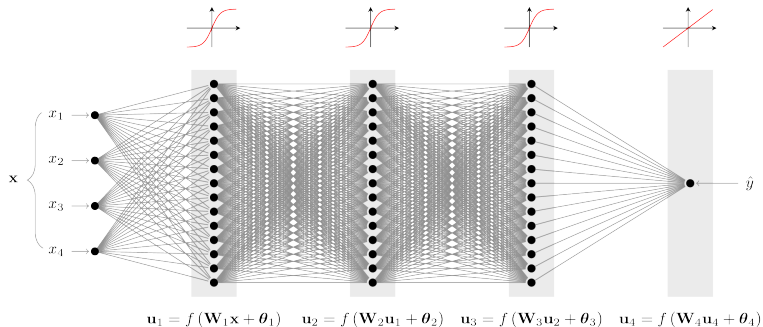
Machine Learning

Neural Networks



Machine Learning

Neural Networks



Machine Learning

Neural Networks

- ▶ (Not so) Recent Trends
 - ▶ Deep networks
 - ▶ End-to-end learning
- ▶ Deep Learning
 - ▶ Complex models, large amounts of data
 - ▶ Enabled through minibatch learning (independence of dataset size) and fast (parallel) CPUs (GPUs)

Machine Learning

Neural Networks Training

- Supervised learning: Minimize mean of loss function $\mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$ over training set $\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^n$.


$$\underset{\mathbf{W}_i, \boldsymbol{\theta}_i}{\text{minimize}} \frac{1}{n} \sum_{i=1}^N \mathcal{L}(F_{NN}(\mathbf{x}_i, \mathbf{W}_i, \boldsymbol{\theta}_i), \mathbf{y}_i) \quad (21)$$

- Use gradient information for efficient training
- Perform training on randomized minibatches (subsets of the training set)

Machine Learning


Neural Networks

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$



Machine Learning

Neural Networks




Machine Learning

Neural Networks

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \theta_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \theta_2)$$




Machine Learning

Neural Networks

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \boldsymbol{\theta}_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \boldsymbol{\theta}_2)$$




Machine Learning

Neural Networks

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \boldsymbol{\theta}_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \boldsymbol{\theta}_2)$$

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$




Machine Learning

Neural Networks

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \boldsymbol{\theta}_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \boldsymbol{\theta}_2)$$




Machine Learning

Neural Networks

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \theta_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \theta_2)$$



$$\frac{dJ_{\mathcal{X}}}{d\mathbf{W}_2} = \left(\mathbf{f}'_2 \odot \frac{dJ_{\mathcal{X}}}{d\mathbf{U}_2} \right)^T \mathbf{U}_1$$

Machine Learning

Neural Networks

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \boldsymbol{\theta}_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \boldsymbol{\theta}_2)$$



$J_X(\mathbf{y}, \hat{\mathbf{y}})$


$$\frac{dJ_X}{d\boldsymbol{\theta}_2} = \left(\mathbf{f}_2' \odot \frac{dJ_X}{d\mathbf{U}_2} \right)^T \mathbf{1}$$

Machine Learning

Neural Networks

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \boldsymbol{\theta}_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \boldsymbol{\theta}_2)$$

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$




$$J_{\mathcal{X}}(\mathbf{y}, \hat{\mathbf{y}})$$

$$\frac{dJ_{\mathcal{X}}}{d\mathbf{U}_1} = \left(\mathbf{f}'_2 \odot \frac{dJ_{\mathcal{X}}}{d\mathbf{U}_2} \right) \mathbf{W}_2$$


Machine Learning

Neural Networks



Machine Learning

Neural Networks





Machine Learning

Neural Networks

$$\mathbf{U}_1 = f_1(\mathbf{X}\mathbf{W}_1^T + \boldsymbol{\theta}_1) \quad \mathbf{U}_2 = f_2(\mathbf{U}_1\mathbf{W}_2^T + \boldsymbol{\theta}_2)$$

$$\begin{bmatrix} x_{0,0} & \dots & x_{0,m} \\ x_{1,0} & \dots & x_{1,m} \\ \vdots & & \vdots \\ x_{n,0} & \dots & x_{n,m} \end{bmatrix}$$



$$J_{\mathcal{X}}(\mathbf{y}, \hat{\mathbf{y}})$$

$$\frac{dJ_{\mathcal{X}}}{d\boldsymbol{\theta}_1} = \left(\mathbf{f}'_1 \odot \frac{dJ_{\mathcal{X}}}{d\mathbf{U}_1} \right)^T \mathbf{1}$$

Machine Learning

Neural Networks

Advantages

- ▶ Computational performance
 - ▶ Optimized CPU/GPU codes readily available
- ▶ Flexibility


Disadvantages

- ▶ Need hyperparameter tuning for optimal performance
- ▶ More-or-less black box models

Machine Learning

Neural Networks Performance


- Intel Xeon Processor E5-1680 v4



Machine Learning

Neural Networks Performance


► NVIDIA Tesla K20 (Single Precision)



Machine Learning

Neural Network Example (Particle Physics)

- ▶ Neural network trained on simulated detector signals (momenta of decay products)
- ▶ Shallow (SNN) and deep (DNN) neural networks
- ▶ Trained with and without hand crafted high-level features (Low, High)



Machine Learning

Neural Networks

- ▶ Advantages:
 - ▶ Simple
 - ▶ Fast
 - ▶ Packages providing optimized code readily available
 - ▶ Flexible
- ▶ Disadvantages:
 - ▶ Need hyperparameter tuning for optimal performance
 - ▶ Black box model

My Current Research

Motivation


- ▶ Neural networks are nice but they usually only yield a single value \mathbf{x} for the retrieval
- ▶ Is it possible to instead retrieve the (approximate) posterior $p(\mathbf{x}|\mathbf{y})$ using a neural network approach?
- ▶ Relevant for the retrieval of (frozen) hydrometeors from passive microwave observations

Approach

- ▶ Learn quantiles of the posterior distribution (quantile regression)

My Current Research

Preliminary Results



References



Christian D. Kummerow et al. “The Evolution of the Goddard Profiling Algorithm to a Fully Parametric Scheme”. In: *Journal of Atmospheric and Oceanic Technology* 32.12 (2015), pp. 2265–2280. DOI: 10.1175/JTECH-D-15-0039.1. eprint: <https://doi.org/10.1175/JTECH-D-15-0039.1>. URL: <https://doi.org/10.1175/JTECH-D-15-0039.1>.